

# CityEar - IoT Urban Noise Monitoring System Report

---

## 1. Introduction

### Problem Statement

Urban noise pollution significantly impacts public health. Monitoring a vast metropolitan area like Hanoi requires a dense, scalable network of sensors continuously collecting data. Traditional approaches using expensive, sparse static stations fail to capture hyper-local events.

### Project Objective

The **CityEar** project designs and safeguards a **Wireless Sensor Network (WSN)** for urban acoustics. We utilize a **Digital Twin** approach to validate the system architecture before physical deployment. The system simulates **1000 independent IoT nodes** transmitting environmental data in real-time, focusing on:

1. **High-Throughput Data Ingestion:** Handling massive concurrent streams.
2. **Low-Latency Anomaly Detection:** instant identification of safety threats (Gunshots, Screams).
3. **Visualization:** Dynamic geospatial mapping of acoustic fields.

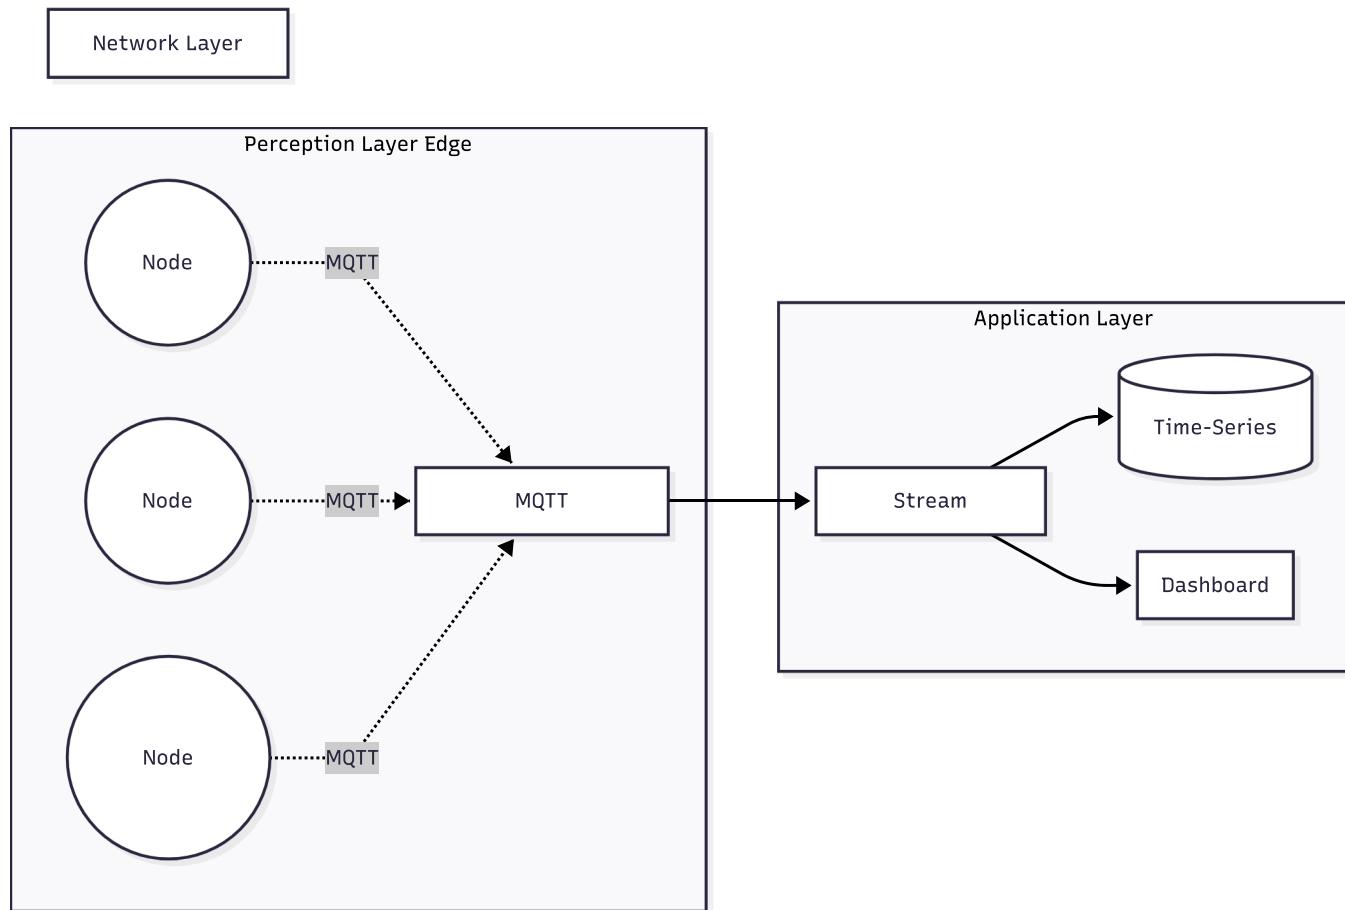
### Team Member

- **Name:** Tran Anh Dung
  - **Student ID:** 20226031
  - **Role:** **IoT System Engineer** (Responsible for Sensor Network Design, Protocol Implementation, and System Integration).
-

## 2. IoT System Design & Analysis

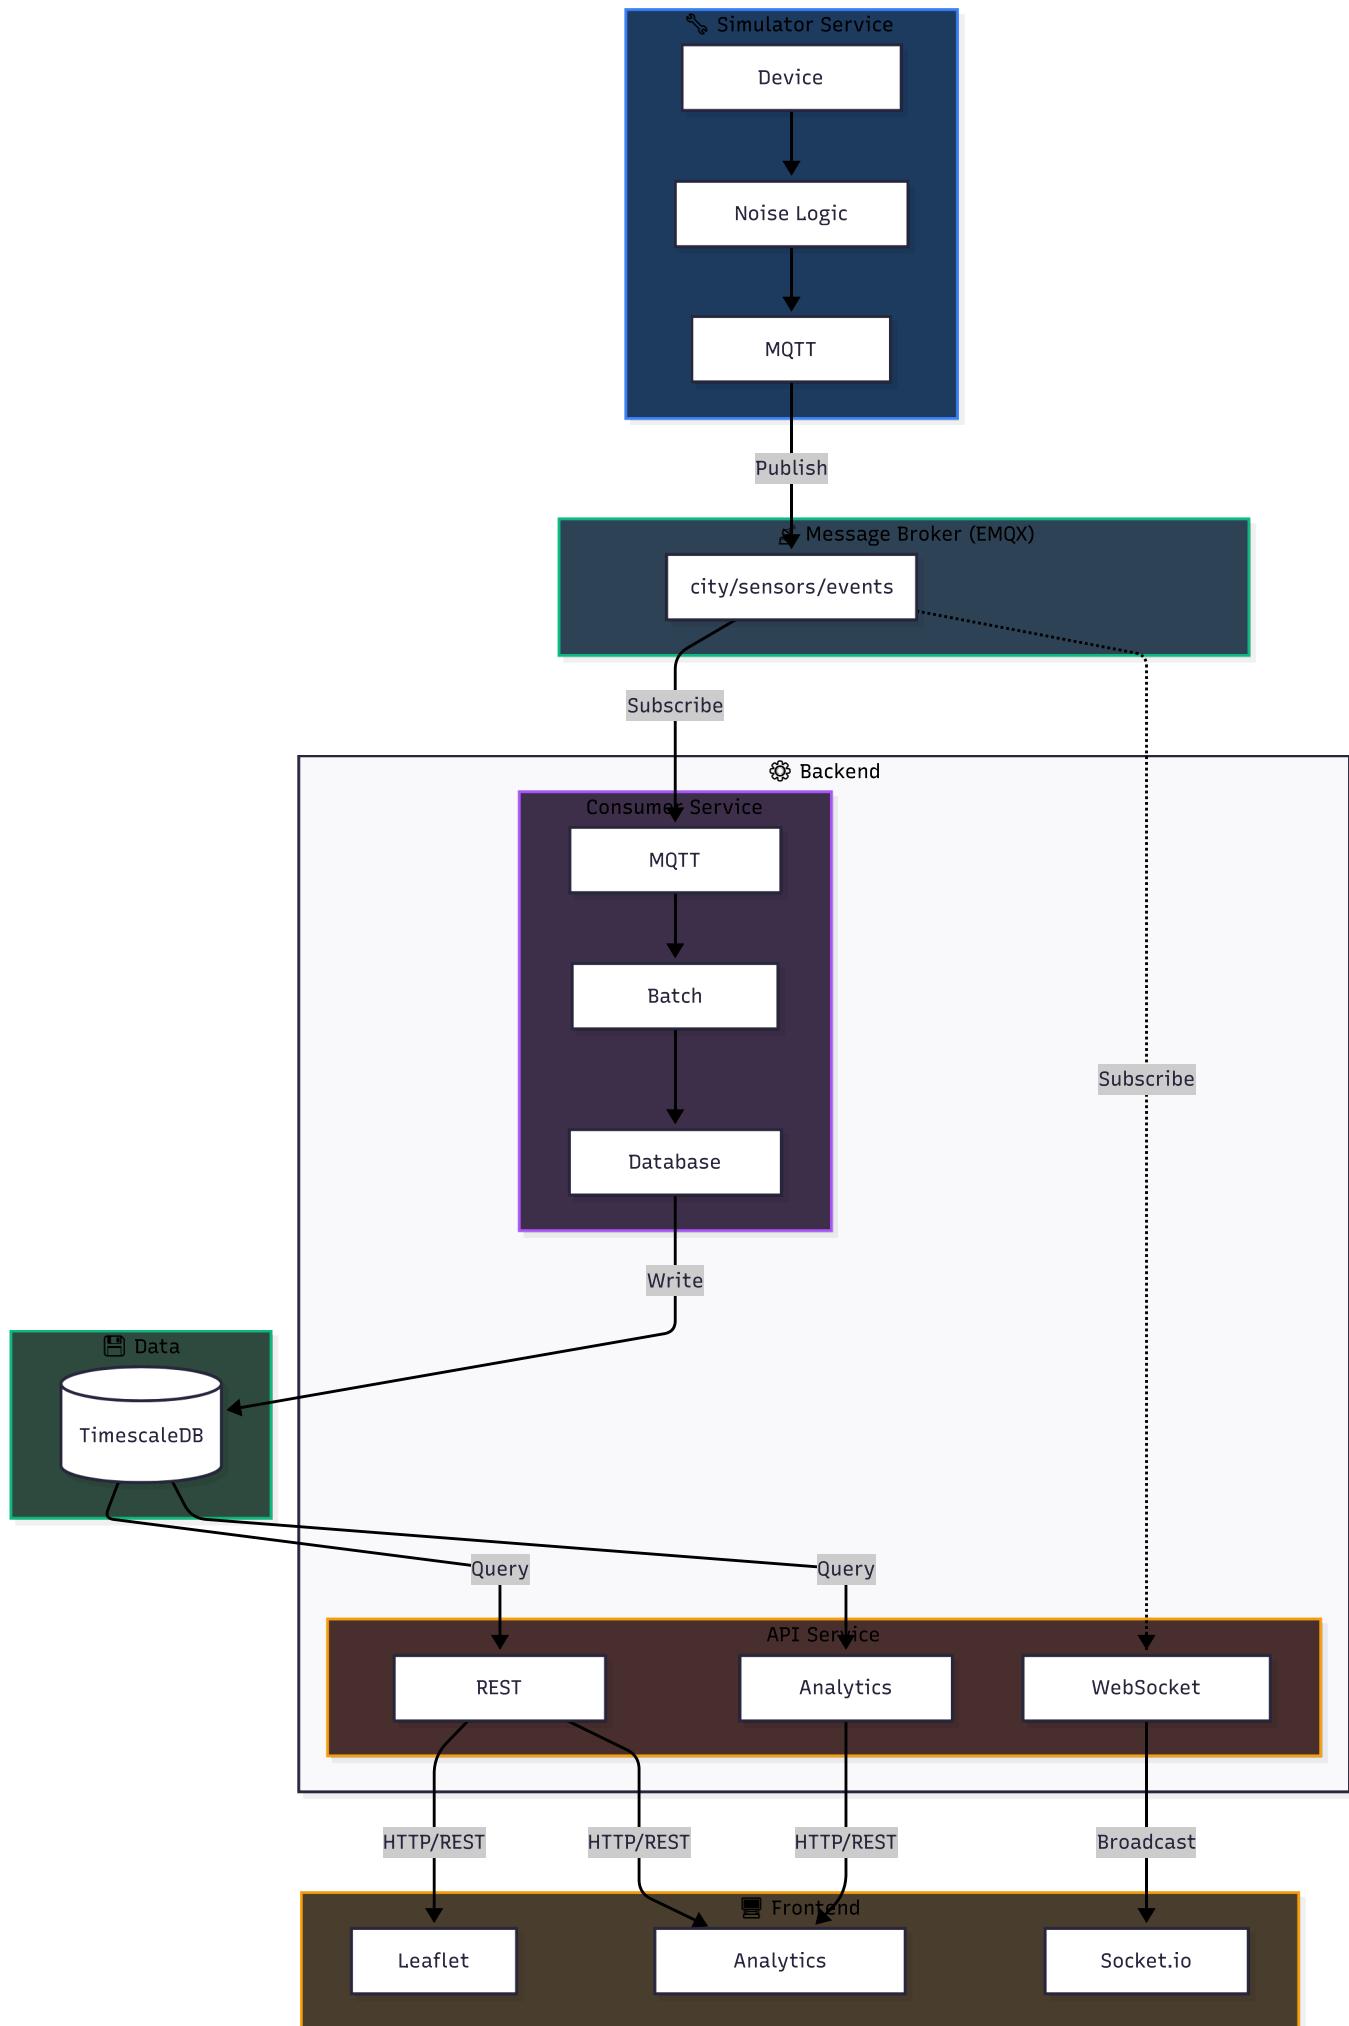
### 2.1. Network Architecture

The system adopts a **Star Network Topology** where individual sensor nodes transmit data directly to a central gateway (Broker). This safeguards against single-point failures in mesh routing and ensures deterministic latency for critical alerts.



### 2.2. Software Architecture & Implementation

The system is built as a **microservices architecture** with clear separation of concerns:



## Service Responsibilities:

### 1. Simulator Service:

- Manages virtual device lifecycle (create, start, stop, delete).
- Generates realistic noise data using time-based state machines.
- Publishes telemetry to `city/sensors/events` topic via MQTT.
- Exposes REST API for device control (port 3001).

### 2. Consumer Service:

- Subscribes to MQTT topics and ingests raw sensor data.
- Implements **batching** to reduce database write pressure.
- Writes events to TimescaleDB hypertables.
- **Critical Path:** Gunshot/Scream events trigger immediate MQTT republish to `city/internal/alerts`.

### 3. API Service:

- **REST API:** Provides HTTP endpoints for historical queries (`/api/events/*`).
- **WebSocket Gateway:** Subscribes to `city/internal/*` MQTT topics and broadcasts real-time updates to connected clients.
- **Analytics:** Aggregates TimescaleDB data for dashboards.

### 4. Frontend:

- **Map View:** Displays device locations with real-time status via Socket.io.
- **Analytics Dashboard:** Fetches aggregated statistics via REST and displays charts.

## Data Flow Example (Gunshot Detection):

1. Simulator generates **GUNSHOT** event → Publishes to `city/sensors/events`.
2. Consumer receives event → Writes to DB → Republishes to `city/internal/alerts`.
3. API Gateway receives from `city/internal/alerts` → Broadcasts via WebSocket.
4. Frontend receives WebSocket event → Updates map marker to pulsing red.

## 2.3. Communication Protocol: MQTT vs HTTP

We selected **MQTT (Message Queuing Telemetry Transport)** over HTTP for the following IoT-specific reasons:

- **Lightweight Overhead:** MQTT headers are small (min 2 bytes), reducing bandwidth usage for power-constrained devices.
- **Pub/Sub Model:** Decouples the sensors from the server, allowing asynchronous communication and easy scaling.
- **QoS (Quality of Service):**
  - **Telemetry (Noise Levels):** Uses **QoS 0** (At most once) – Occasional packet loss is acceptable for continuous streams.
  - **Alerts (Gunshots):** Uses **QoS 1** (At least once) – Critical safety data must be guaranteed to arrive.

## 2.3. Data Model & Payload Optimization

To minimize radio-on time (power consumption), payloads are kept minimal JSON.

### Telemetry Packet (Periodic):

```
{
  "id": "dev_01",
  "v": 65.4,    // Noise Level (Float)
  "ts": 1712345678
}
```

### Alert Packet (Event-Driven):

```
{
  "id": "dev_01",
  "t": "GUNSHOT", // Event Type
  "v": 125.0,
  "lat": 21.02,
  "lng": 105.85,
  "ts": 1712345678
}
```

## 3. Component Implementation

### 3.1. Perception Layer (The Simulator)

Instead of physical hardware, we implemented a **High-Fidelity Software Simulator** using **Node.js (NestJS)** to model the behavior of a city-wide sensor network. This allows us to stress-test the backend without the cost of deploying 1000 sensors.

- **Virtual Device Design:**
  - Each "Device" is an independent object with a unique ID and fixed geolocation (jittered around real Hanoi coordinates).
  - **Time-Aware Noise Logic:** The simulator implements a state machine based on the time of day:
    - **Rush Hours** (07:00-09:00, 17:00-19:00): Base noise set to **80dB** to simulate heavy traffic.
    - **Night Hours** (23:00-05:00): Base noise drops to **40dB** (Ambient/Quiet).
    - **Normal Hours:** Base noise averages **60dB**.
  - **Randomness & Jitter:** Statistical variance (+/- 5dB) is applied to every reading to prevent artificial "flatline" data.
- **Anomaly Injection (Probabilistic Model):**
  - To test the critical alert pipeline, the simulator injects synthetic acoustic events based on probability:
  - **Gunshots (0.5% chance):** Spikes to **120-130dB**.
  - **Screams (0.5% chance):** Spikes to **95-105dB**.
- **Transmission:** Devices publish JSON payloads to the MQTT broker on the **city/sensors/events** topic every 5 seconds.

## Device Management Interface:

The simulator provides a web-based control panel for managing virtual devices:

The screenshot shows a dark-themed web application titled "Device Management" with the subtitle "Control Plane for Sensor Network Simulation". At the top right is a green button labeled "+ Deploy Sensor". Below the header is a search bar with placeholder text "Search devices...". To the right of the search bar is a green button with the text "500 Active Sensors". The main area is a table with the following columns: ENVIRONMENT, PROFILE, STATUS, and LOCATION. The table lists 10 rows of sensor data, each with a small icon representing the environment (HO or CA), the device ID (e.g., device-0000 to device-0010), the profile (QUIET\_RESIDENTIAL), the status (Online with a green dot), and the location coordinates (e.g., 21.0274, 105.8556). At the bottom left is a "Show" dropdown set to "10" with a downward arrow. At the bottom right is a page navigation bar showing "Page 1 of 50" with left and right arrows.

Figure 1: Device management interface showing active/inactive devices with real-time status

The screenshot shows a dark-themed "Add New Device" form. At the top right is a close button "X". Below it is a sub-header "Deploy a new sensor to the network.". The form has several input fields: "Device Name" with placeholder "e.g., Hoan Kiem Sensor #9000", "Latitude" set to "21.0285", "Longitude" set to "105.8542", "Profile" set to "Quiet Residential", and an "Active Status" toggle switch which is turned on. At the bottom right are two buttons: "Cancel" and a green "Create Device" button.

Figure 2: Device creation form with configurable noise profiles

## 3.2. Network & Processing Layer

- **Broker (EMQX):** Handles 1000 concurrent TCP connections.
- **Stream Processing:**

- **Batching Strategy:** Non-critical data is buffered (Windowing) and written to disk every 1 second to reduce database I/O pressure.
- **Fast Path:** "Gunshot" events bypass the buffer and trigger immediate WebSocket broadcasts.

### 3.3. Application Layer (Visualization)

- **Heatmap:** Aggregates spatial data to show noise pollution "hotspots" rather than individual raw points.
- **Real-time Alerts:** Provides sub-second visual feedback for emergency response.

#### Real-time Noise Map:

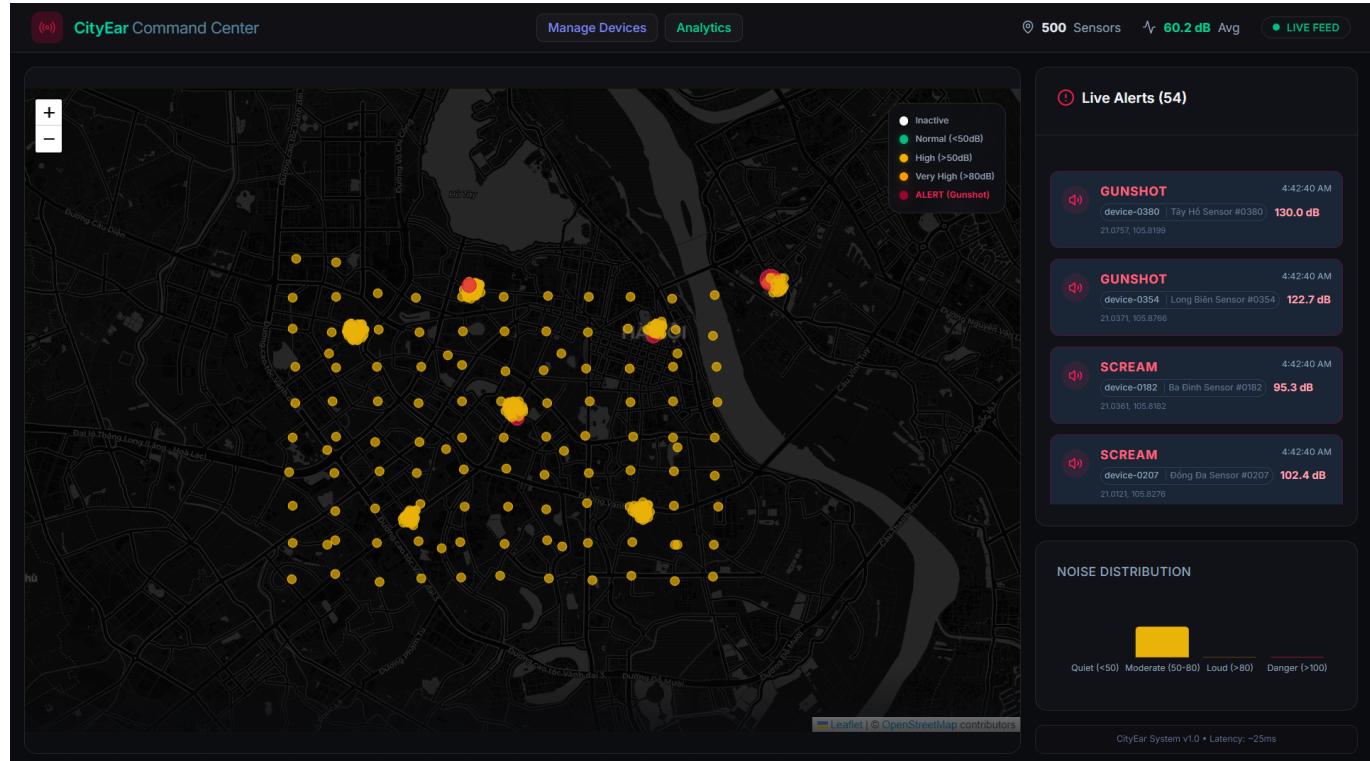


Figure 3: Interactive map displaying sensor locations with color-coded noise levels and anomaly alerts

#### Analytics Dashboard:

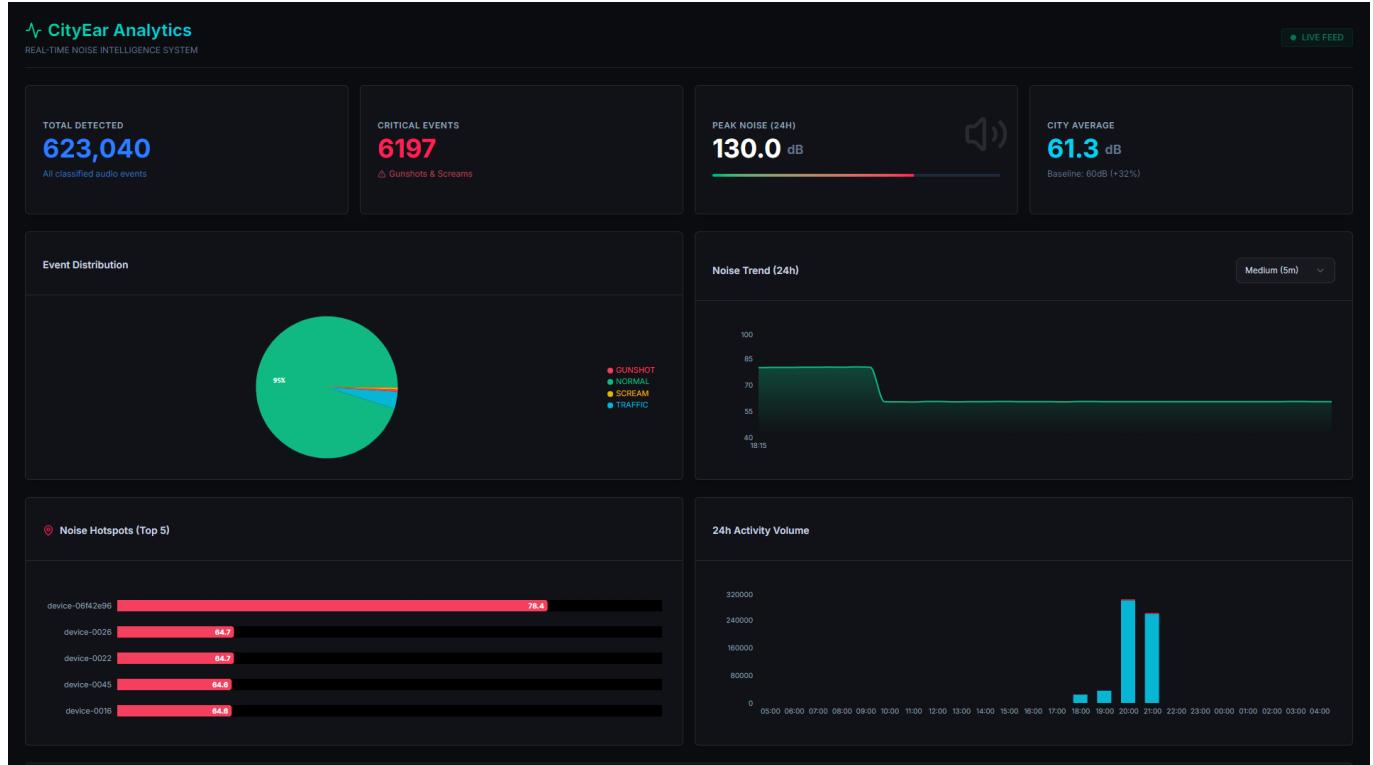


Figure 4: Analytics dashboard showing noise trends, event distribution, and top noise hotspots

## 4. Experimental Results

### 4.1. Performance Metrics

- Simulated Load:** 1000 Devices.
- Message Rate:** ~200 messages/second.
- End-to-End Latency:**
  - Telemetry Update: ~300ms.
  - Critical Alert: < 150ms (due to "Fast Path" routing).

### 4.2. Scalability

The use of **TimescaleDB** (Hypertable) allows the system to ingest millions of rows per day without query performance degradation, proving the architecture is suitable for city-wide deployment.

## 5. Future Work

- Hardware Deployment:** Replace simulator nodes with physical ESP32 prototypes.
- Edge AI:** Implement TinyML on the ESP32 to classify audio sounds (e.g., distinguishing a jackhammer from a gunshot) locally before transmission.
- Low-Power WAN:** Investigate LoRaWAN for wider coverage in areas without Wi-Fi.

## Appendix: System Screenshots

All figures referenced in this report demonstrate a fully functional prototype system handling 500+ simulated IoT devices with real-time visualization and analytics capabilities.